

Gleb Chuvpilo
Knight Capital
Order Book Problem
November 15, 2012

Running the solution using a market data file:

```
$ make clean && make && make run
```

Running the solution using standard input (output in bold):

```
$/Pricer 200
28800538 A b S 44.26 100
28800562 A c B 44.10 100
28800744 R b 100
28800758 A d B 44.18 157
28800758 S 8832.56
28800773 A e S 44.38 100
28800796 R d 157
28800796 S NA
28800812 A f B 44.18 157
28800812 S 8832.56
28800974 A g S 44.27 100
28800974 B 8865.00
28800975 R e 100
28800975 B NA
28812071 R f 100
28812071 S NA
28813129 A h B 43.68 50
```

```
28813129 S 8806.50
28813300 R f 57
28813300 S NA
28813830 A i S 44.18 100
28813830 B 8845.00
28814087 A j S 44.18 1000
28814087 B 8836.00
28814834 R c 100
28814864 A k B 44.09 100
28815774 R k 100
28815804 A l B 44.07 175
28815804 S 8804.25
28815937 R j 1000
28815937 B 8845.00
28816245 A m S 44.22 100
28816245 B 8840.00
```

Files/line count:

```
$wc -l *.h *.cpp
    89 Exceptions.h
   186 Log.h
    75 MarketDataProvider.h
    58 MarketOrder.h
   167 OrderBook.h
    87 Parser.h
    42 Utils.h
   168 Main.cpp
    30 MarketDataProvider.cpp
   109 OrderBook.cpp
    41 Utils.cpp
  1052 total
```

What is the time complexity for processing an Add Order message?

I am using an STL multimap to store orders using price as the key. STL multimap is implemented as a binary search tree. Therefore, Add Order messages are processed $O(\log n)$.

What is the time complexity for processing a Reduce Order message?

I am using an additional data structure – STL unordered_map – to go from an order ID to an iterator into the multimap above. The STL unordered_map is implemented as a hash map, therefore lookup is $O(1)$. I pass this iterator to the multimap's erase method, which runs $O(1)$. Thus, total time complexity of Reduce Order is $O(1)$.

If your implementation were put into production and found to be too slow, what ideas would you try out to improve its performance?

- Increase the size of the hash
- Improve the hashing function
- Use a custom multimap implementation
- Use a cache-aware algorithm
- Inline more methods (most are inlined already, though)
- Use better message parsing
- Use an GPU or an FPGA (preferable)
- Co-locate
- Use a faster/fatter optical link to the exchange
- Use better network cards and routers
- Timestamp packets on arrival and on departure and figure out the bottleneck
- Use valgrind for memory optimization